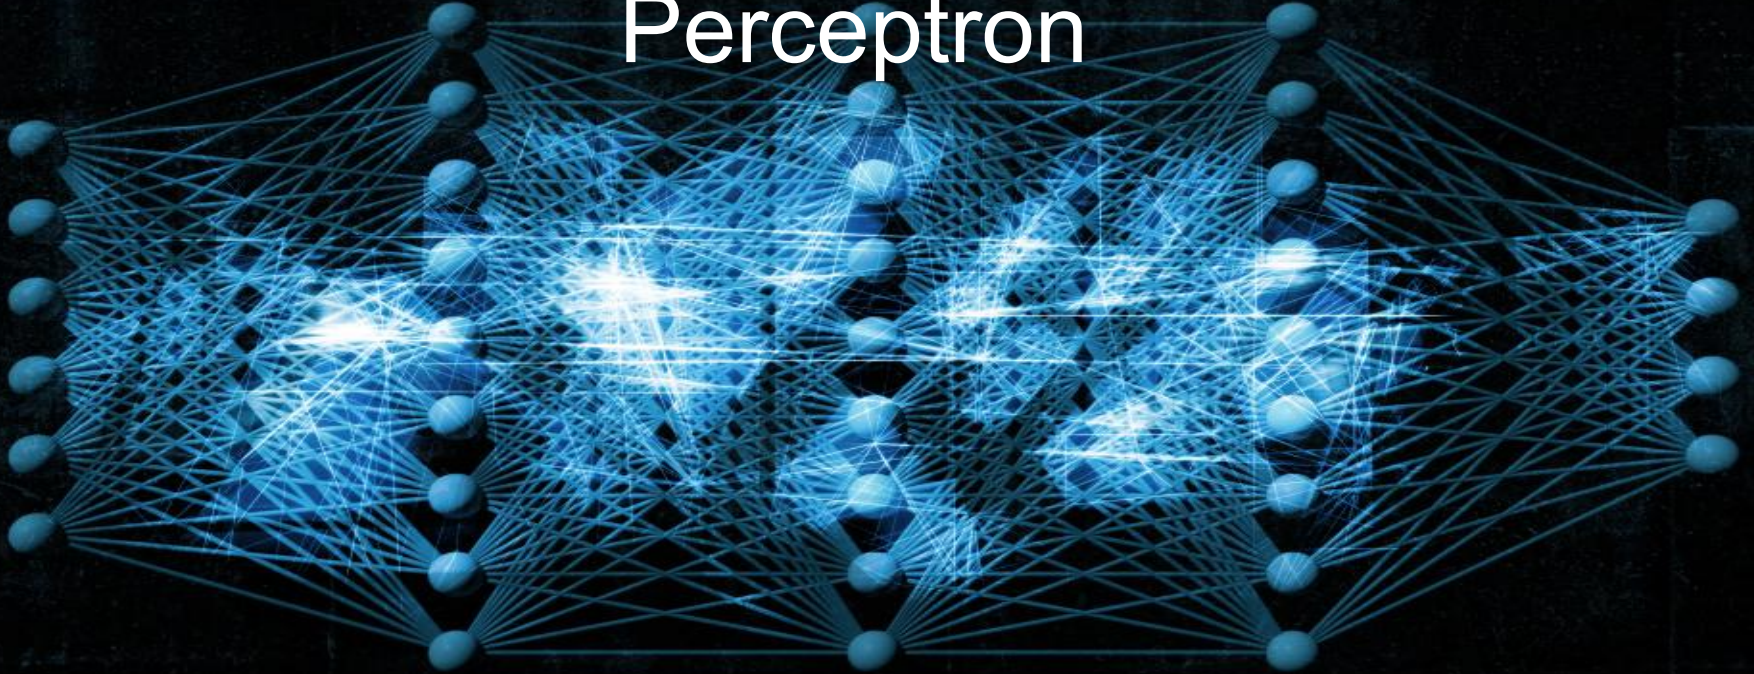


CUDA Accelerated Training of Multilayer Perceptron



Dataset Description

The CIFAR-10 dataset is a collection of images that are commonly used to train machine learning and computer vision algorithms. The dataset consists of:

- ❑ 60000 32x32 colour images in 10 classes,
- ❑ with 6000 images per class.
- ❑ There are 50000 training images and 10000 test images

airplane



automobile



bird



cat



deer



dog



frog



horse



ship

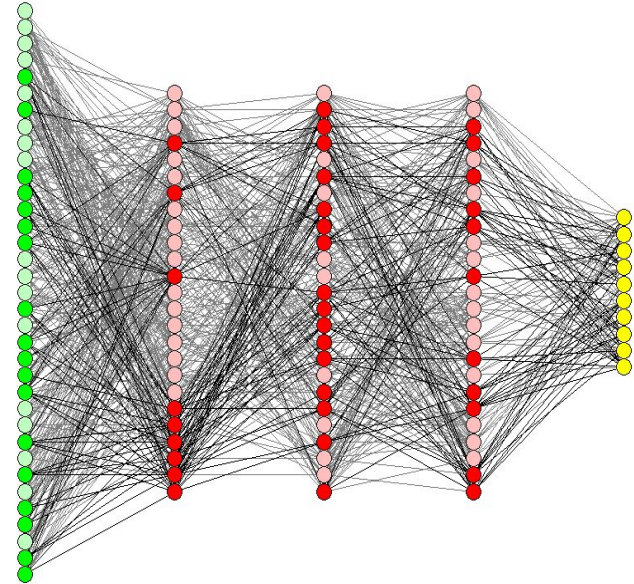


truck



The Multilayer Perceptron Architecture

- ❑ Input layer : 3072 ($32 * 32 * 3$) nodes
- ❑ First Hidden Layer: 2048 nodes
- ❑ Second Hidden layer: 1024 nodes
- ❑ Third Hidden layer: 512 nodes
- ❑ Output layer: 10 nodes (numbers of classes)



Matrices Operations Kernels:

We implement kernels for matrices operations in the neural network layers

- Dot_Product_Kernel
- Multiplication_Kernel
- Addition_Kernel
- Subtraction_Kernel
- ReLU_Activation_Kernel
- Derivative_ReLU_Kernel
- Tanh_Activation_Kernel
- Exponentiation_Kernel
- Power_Kernel
- Sum_Kernel
- Division_Kernel

See matrix.cu file for detail implementations

```
-global void Dot_Product_Kernel(Matrix *A, Matrix *B, Matrix *C, bool trans1, bool trans2) {
    double Cvalue = 0.0;
    int row = threadIdx.y + blockIdx.y * blockDim.y;
    int col = threadIdx.x + blockIdx.x * blockDim.x;
    if (trans1 == false && trans2 == false)
        if (row < A->height && col < B->width) {
            for (int i = 0; i < A->width; i++)
                Cvalue += Get_Mat_element(A, row, i) * Get_Mat_element(B, i, col);
            Set_Mat_element(C, row, col, Cvalue);
        }
    if (trans1 == true)
        if (row < A->width && col < B->width) {
            for (int i = 0; i < A->height; i++)
                Cvalue += Get_Mat_element(A, i, row) * Get_Mat_element(B, i, col);
            Set_Mat_element(C, row, col, Cvalue);
        }
    if (trans2 == true)
        if (row < A->height && col < B->height) {
            for (int i = 0; i < A->width; i++)
                Cvalue += Get_Mat_element(A, row, i) * Get_Mat_element(B, col, i);
            Set_Mat_element(C, row, col, Cvalue);
        }
}

-global void Multiplication_Kernel(Matrix *A, Matrix *B, Matrix *C) {
    double Cvalue = 0.0;
    int row = threadIdx.y + blockIdx.y * blockDim.y;
    int col = threadIdx.x + blockIdx.x * blockDim.x;
    if (row < A->height && col < A->width) {
        Cvalue = Get_Mat_element(A, row, col) * Get_Mat_element(B, row, col);
        Set_Mat_element(C, row, col, Cvalue);
    }
}

-global void Multiplication_Kernel(Matrix *A, double k) {
    double Avalue = 0.0;
    int row = threadIdx.y + blockIdx.y * blockDim.y;
    int col = threadIdx.x + blockIdx.x * blockDim.x;
    if (row < A->height && col < A->width) {
        Avalue = Get_Mat_element(A, row, col) * k;
        Set_Mat_element(A, row, col, Avalue);
    }
}

-global void Addition_Kernel(Matrix *A, Matrix *B, Matrix *C) {
    double Cvalue = 0.0;
    int row = threadIdx.y + blockIdx.y * blockDim.y;
    int col = threadIdx.x + blockIdx.x * blockDim.x;
    if (B->height == 1)
        if (row < A->height && col < A->width) {
            Cvalue = Get_Mat_element(A, row, col) + Get_Mat_element(B, 0, col);
            Set_Mat_element(C, row, col, Cvalue);
        }
}

-global void Subtraction_Kernel(Matrix *A, Matrix *B, Matrix *C) {
    double Cvalue = 0.0;
    int row = threadIdx.y + blockIdx.y * blockDim.y;
    int col = threadIdx.x + blockIdx.x * blockDim.x;
    if (row < A->height && col < A->width) {
        Cvalue = Get_Mat_element(A, row, col) - Get_Mat_element(B, row, col);
        Set_Mat_element(C, row, col, Cvalue);
    }
}
```

MLP Training Functions:

- `prepare_data()`: to prepare the dataset for the training
- `predict_y()`: to predict the output of the model
- `calculate_loss()`: to calculate the loss of the neural network on a batch of data
- `forward_propagation()`: to pass the data through the model architecture
- `backward_propagation()`: for updating the weights.

See main.cu file for detail implementations

MLP Implementation in Pytorch

We implemented the same architecture with:

- the same hyperparameters
(lr = 0.0001, Batch size = 64),
- optimization algorithm
(Stochastic Gradient Descent)
and
- same loss function
(Cross entropy loss)



Results and Comparison